

Vittorio Rossi

AI Engineer & Researcher

rossi.vittori@gmail.com | +39 3458367454 | vittoriorossi.com | linkedin.com/in/vitrossi | github.com/VittorioRossi | Google Scholar

PROFESSIONAL EXPERIENCE

Commerce Clarity <i>Founding AI Engineer</i>	June 2024 – Present Milan, Italy
- Joined pre-team as the first AI hire; built and own the workflow execution engine running the company's product-enrichment, translation and data-sourcing agent pipelines - Built the LLM infrastructure layer for prompt versioning, multi-provider clients and execution analytics, tracking 10M+ production calls - Fine-tuned a Gemma 3 model with LoRA on a MacBook to adapt tone of voice for a client - Built the pipeline generating a 1M-product catalog, plus a vision pipeline matching 10K products against 5M competitor listings - Defined the company's evaluation standard: a standardized eval framework setting out what evaluation means in this context, with rubrics and a fixed spec format reused across customer workflows - Built the evaluation harnesses behind it: an LLM-as-judge scorer and a general harness running the framework across customer workflows	
Q8 Kuwait Petroleum Italy S.p.A. <i>Data Scientist Intern</i>	June 2023 – November 2023 Rome, Italy
Led a 4-person team redesigning the IoT backend ingesting video streams from 200+ IoT devices into millions of data points, reducing measurement errors by 50%; migrated 5+ pipelines to Apache Spark, AWS Glue, and AWS Redshift, triangulating device telemetry against SQL-warehoused data	

TECHNICAL PROJECTS

IRE (Integrated Research Environment) <i>Co-founder</i>	May 2026 – Present
Co-founded and released an open-source, local-first research environment for ML researchers, built in Rust and TypeScript, with macOS and Linux releases under the MIT license - Designed a pluggable agent-provider interface unifying Claude Code, Codex and OpenCode, and a native MCP bridge giving agents a Git-tracked wiki, structured memory and non-blocking experiment execution	
GraphCode-Bench <i>LLM Evaluation Benchmark</i>	March 2026
Built a benchmark for agentic code understanding over graph-structured repository representations, comparing five open-weight models: Qwen3 14B, Qwen3-Coder 30B, GPT-OSS 20B, Mistral-Small 24B, and Llama-3.1 8B	

RESEARCH & PUBLICATIONS

Diffusion Language Models Are Natively Length-Aware	Accepted, NeurIPS 2026 Workshop on Diffusion Language Models
<i>Vittorio Rossi, Giacomo Cirò, Davide Beltrame, Luca Gandolfi, Paul Röttger, Dirk Hovy</i>	
First author. Showed that masked diffusion LLMs encode the required output length before decoding, and turned that into a zero-shot, training-free method shortening the generation canvas from a single forward pass >30% FLOPs reduction on every benchmark, up to 90% on some, with no statistically significant performance loss, and improved performance on 2 of 4 (GSM8K, HumanEval, IfEval, LongFormQA on LLaDA-8B) - Ran the experimental program on H100 GPUs via SLURM; owned hypothesis through manuscript	

EDUCATION

Bocconi University <i>MSc in Artificial Intelligence, exam average 28/30</i>	September 2024 – December 2026 Milan, Italy
Georgia Institute of Technology <i>Exchange Programme</i>	January 2024 – May 2024 Atlanta, USA
Bocconi University <i>BSc in Economics, Management and Computer Science, 110/110</i>	September 2021 – July 2024 Milan, Italy

TECHNICAL SKILLS

Languages	Python (expert, Pandas, NumPy), TypeScript, Rust, SQL, PHP, C/C++
ML and LLM Systems	PyTorch, Transformers, LoRA, diffusion LMs, inference optimisation, evaluation and benchmarking
Infrastructure and Data	SLURM and HPC (H100), Kubernetes, Docker, Google Cloud Platform, CI/CD, Neo4j, MongoDB
Spoken	Italian (native), English (fluent, IELTS 7.5), Spanish (basic)